AI ETHICS

AI Governance and AI in Government

By Smitha Krishna Prasad, Fritz Family Fellow, Georgetown University Law Center

Today, governments are turning to the use of AI to improve their 'GovTech', the technology that is developed and used for their own functions and services. However, most policy and governance discussions on AI stick to questions about how big tech and the private sector's development and use of AI must be governed. This module provides an overview of some ways governments are using AI in their own functions, the policies that they follow in this context, and some of the issues and concerns that come up with extensive use of AI in government, and particularly for citizen services.

Lecture Transcript

0:02 Hi, everyone. I'm Smitha Krishna Prasad. I'm a lawyer working on issues of technology, law and policy, particularly in the context of privacy, data protection and GovTech. That is the technology used for government functions. I'm currently a Fritz Family Fellow with Georgetown's Tech and Society program and also an Assistant Professor of Law at the National Law School (of India) University, Bangalore. Prior to this, I've spent several years working on emerging tech policy issues in India and globally, and have led work on these issues as a Director at the Center for Communication Governance at the National Law University, Delhi. Currently, my work focuses on policy and legal framework applicable in relation to the use of digital technologies and government services, and the development of public digital infrastructure at scale in India and in the US. I'm also interested in how these issues play out across the Global South countries and the role that the international organizations play in this context.

1:06 So while you generally hear a lot about ethics and regulation of AI, these conversations typically focus on Big Tech companies and more broadly, the private sector's engagement with AI. While efforts to regulate the way companies build and operate AI are critical, I wanted to talk about another important aspect of AI, which is just as, if not more, important. And that's the governance policies that we have on the way in which governments themselves use these technologies. So while most countries have been going through the process of digitizing and modernizing their administrative offices for some time, we see that today, there's also a push that is, in some cases pretty aggressive, to incorporate the use of AI in these processes. This push often comes in as part of the overall policies the government may have on the promotion, development and use of AI in the country. But in addition to ensuring that the domestic industry is competitive in the space, there's also usually an interest in harnessing the benefits of AI towards providing more value to citizens. And one of the ways you do that is by making governance itself more efficient and access to government more efficient.

2:25 As we debate the kind of regulation that is needed to ensure that the private use of AI doesn't violate individual human rights, we also need to consider what parameters may be required when governments use such technology since there's so much that's happening. Over the next few minutes or so, we will go over the different ways in which governments are engaging with AI as it's used within government bodies. To understand these policies better, we'll also look at some examples of some of the jurisdictions that have been either most vocally engaged in this debate, whether in terms of policy or practice, and others also that civil society actors, especially I'd argue, should be brought into the debate in a more local way, because of developments that are happening in these countries.

3:12 So like I said, while several countries have policies and regulations that directly and indirectly address AI, these policies are often, you know, in the nature of broader industrial policies that focus on the promotion and development of AI, and the promotion and development of the AI industry in the domestic economy. However, obviously, policy is not just enough, and we look beyond policy to regulation to see how these policies are implemented, and how this would come into effect. And in this context, you may have heard in the news that the EU very recently agreed on the AI Act, which specifically targets and regulates AI in the past few months. But often in other jurisdictions, we need to rely on a mix of horizontally applicable laws. So for instance, privacy law, which applies across all sectors, and then sector specific laws and regulations that kind of impact the specific industry, or the specific companies that we're looking at. And this is how you understand what is permitted and what is not, whether in the context of AI or otherwise.

4:25 In addition to policy and regulation, we also see there are other actors in the space that includes: courts, and international coalitions or international organizations, civil society actors, as well as industry stakeholders themselves. So we've seen over the past few years that several sets of AI principles and ethics have emerged from the work of all of these different stakeholders and organizations. But as I mentioned earlier, one issue that's not always discussed in detail in this context is how governments use AI. Policies will typically have sections focused on public sector use of AI or the ideal of using AI for the public good. However, details about the practical implementation of these policies and the mechanisms that exist for accountability and regulation, or should exist for this, are often lacking. So we'll now look at what's happening in practice and in policy in a few jurisdictions. Specifically, we're going to look at India, the US, Singapore and Jordan. And now I recognize that many other countries, particularly if you think about, you know, places like Estonia, Taiwan or China, they're known to have huge investments made into the use of AI for the public sector as well. And we also know that with the support of multilateral organizations, such as the World Bank or the UNDP, several lower and middle income countries are also putting automated systems in place. This is often in relation to their welfare distribution systems. And we'll see some examples of that. So I just wanted to mention that while there are several countries that do this, the ones that we are going to discuss here are examples which are important for different reasons, but they also might not be representative of all countries' engagements with AI in government functions.

6:22 So first, we'll look at how India uses AI in its GovTech programs or in its government administration programs. India has several sets of GovTech programs running across both central and state government agencies. As these programs develop, government agencies are incorporating AI in a variety of services and applications. One of the most common uses is, of course, facial recognition technology. For instance, in India, you see that facial and biometric recognition based technologies have been incorporated into what the government calls the DigiYatra scheme, which is used in the context of airport security. But we also see many instances where there are plans to incorporate facial recognition technology into law enforcement and policing activities. On the other hand, if you look at corporate and financial regulators, we see that the Ministry of Corporate Affairs is looking to use data analytics and machine learning to simplify regulatory filings. The securities regulator, SEBI, has suggested that the use of AI will help them put together more tailored regulations based on risk levels seen within different entities. And then when we look at kind of how the government engages with citizens, more specifically, a big example of AI use in public engagement is the use of AI chatbots and natural language processing technologies, which helps translate between the many languages spoken in the country and ideally makes government more accessible to the people who are using these technologies. We’ve also seen, as I mentioned, that the Indian government has several existing GovTech programs, which includes Aadhaar, which is the digital ID system that India is known for, as well as a retail payment system called the UPI. And the government has announced that AI applications will be layered into these existing programs.

8:34 And then other more controversial uses that have been announced and not really materialized. For instance, over the years the Indian government has put out plans and tenders for technology to monitor social media posts made by Indian citizens, but a lot of them were withdrawn. And last, we see that the Indian government has suggested that in January 2024, they will make some big announcements about future plans in this space. So let's look at the context in which we must place these developments. So there's a lot of you know, policies and statements that call for harnessing AI to further government processes in India. And you know, some of these statements and policies set out principles and frameworks for this purpose. Of course, like I mentioned, these are not, you know, specific to government use of AI but more general as well sometimes. So, India has two sets of such documents. The first is the National Strategy for AI, which is a discussion paper from 2018 that focused on five sectors that are envisioned to benefit the most from AI in serving societal needs. These sectors were: healthcare, agriculture, education, Smart Cities and infrastructure, and mobility and transportation. This strategy talks about the importance of building a robust AI ecosystem in India, looking at both uses for the public sector and benefits for the private sector.

10:16 So the second set of documents or policies is a discussion paper series on responsible AI, and responsible AI for all specifically, which was published by the Niti Aayog, which is a think tank and resource center within the government. So, the first of these three discussion papers outlines the principles needed in this context, you know, to kind of build out and develop responsible AI for all. These principles are: safety and reliability, equality, inclusion and non discrimination, privacy and security, transparency, accountability, and protection and reinforcement of positive human values. The second document puts together an enforcement framework, which focuses on the actions that can and should be taken by governments in terms of designing regulatory and policy interventions. It recommends a risk based system which focuses on both protection of individual rights and efforts to remove barriers in the adoption of AI. It also suggests an independent multidisciplinary advisory group at the central level, a council for ethics and technology specifically, which will have the authority to cover the entire digital sector and serve as a think tank, advising different sectoral regulators and engage in research on AI governance and ethics. And then the third discussion paper here looks at applying these frameworks to a specific use case and that's the use of facial recognition technology.

12:03 There are other efforts that we've seen from the Indian government. They have launched the National Program on AI as an umbrella program, which seeks to leverage transformative technology to foster inclusion, innovation and adoption for social impact. There’s a national AI portal, which is supposed to be a repository of AI based initiatives in the country. And as I mentioned earlier, there are announcements about how 2024 will also see many more developments in this space. In terms of regulation, however, there is no specific regulation that addresses AI in India. Some of the laws and regulations that could impact AI include the new Digital Personal Data Protection Act, as well as more narrow rules that exist on you know, social media content moderation, a newer rule on dark patterns and the consumer protection, however most of these laws and regulations are targeted towards the private sector, and do not apply to government actions.

13:15 Similarly, when we look at the US, we see that there are several instances of, you know, federal agencies using AI in their GovTech systems, as well as state agencies. A study by academic researchers at Stanford and NYU shows that over 45% of US federal agencies have used automated systems or tried to use them. However, they also suggest that only 12% of these systems were highly sophisticated uses of automated systems. While many of the more well known examples of automated systems and government involve either welfare or public distribution systems, including the social security systems or facial recognition in law enforcement, this study showed that there are multiple use cases for AI in government and classifies these into five categories, which include: enforcement; regulatory research, analysis and monitoring; adjudication; public services and engagement; and internal management. When we look to its policy, in the US, over the past few months, there's been some significant developments on AI related policy. As you may have heard, the White House issued the AI Executive Order, which covers both the need for promotion and regulation of private sector use of AI and the promotion and regulation of public sector use of AI. Specifically, there's an entire section that addresses federal agency use of AI and it provides that every agency must appoint a Chief AI Officer and study the potential use cases for AI and promote AI innovation within the agency.

15:13 Now, subsequent to the executive order, we've also seen that the Office of Management and Budget published a draft policy on advancing governance, innovation and risk management for agency use of artificial intelligence. Now, this draft emphasizes steps to both remove barriers to the use of AI in agencies and also to manage the risks that arise from the use of AI, and could potentially pave the way for, you know, a framework or a regulation in the space. However, when we look towards law and regulation right now, there's some disparate sets of laws that could apply in this context. For instance, that is something called the AI in Government Act. But the main purpose of this act is to improve the use of AI across the federal government by providing access to technical expertise and streamlining hiring within the agencies. So it doesn't really talk about how they should be using AI and what kind of, you know, restrictions or parameters that it should operate within. However, as a more clear example of an issue or rights based restriction, we also see that a number of jurisdictions, you know, cities and states and local governments in the US have made efforts to ban the use of facial recognition systems, particularly in the context of law enforcement use. And we also see that existing regulations such as the federal Privacy Act may also address some issues, but it's probably not quite enough at a federal level.

17:02 Moving to Singapore. Here, we see that, you know, GovTech and AI is much more cross cutting across the government. The Singapore government, for instance, has a chatbot called OneService, which allows citizens to lodge complaints or request information on them using, you know, social messaging apps like WhatsApp and Telegram. This chatbot is powered by AI and it can automatically identify and classify complaints into appropriate categories, extract the details of the incident that needs attention, and identify and inform the relevant government agency to follow up on the case. So citizens do not really need to figure out which agency or who they need to contact, they can just put that information into this chatbot and the rest of the work will be done. Other examples of AI and public services include adaptive learning systems in schools, chronic health management systems in hospitals, AI uses to support immigration and customs clearance, as well as uses to detect and deter online scams.

18:30 All of these technologies, you know, come up in the context of the national AI strategies that Singapore has. While there was originally a strategy from 2019, in the past month, National AI Strategy 2.0 has been published. This new strategy has a vision of AI for the public good for Singapore and the world. It focuses on twin goals of excellence and empowerment to achieve this vision. And one of the aims of the strategy is to work within government as one of three stakeholders, the other two being corporates and researchers, so to work within government to enhance public sector use of AI. And this includes a drive to improve public service productivity through the use of both general and specialized AI services. We also see that the Personal Data Protection Commission of Singapore has released two editions of a model AI governance framework, along with an implementation and self assessment guide for organizations. So not quite a regulation, but some, you know, involvement by regulators here.

19:46 So the last example we'll look at is Jordan. Jordan is interesting because it's an example of a very different context for the use of automated systems and government. The particular system that we are looking at is something that has been promoted by the World Bank and has become controversial because it, you know, potentially harms the people that it is intended to support and provide for. The system itself is a Unified Cash Transfer program, and the automated system that is used in this context helps decide who should receive cash transfers based on economic vulnerability rankings. And I bring this example up to highlight the role that multilateral organizations like the World Bank play in this space. There are other organizations including the IMF, UNICEF, the United Nations Development Programme and the WFP that are funding the use of automated systems by governments, particularly for welfare delivery or public services, such as the example we see in Jordan. So now, while the stated objectives for these systems are typically to use, you know, data based evidence or data intensive systems to increase efficiency and effectiveness in operating these schemes that provide the targeted assistance. These systems are typically quite controversial, because many reports suggest that they're not designed with an understanding of the people at the other end and the circumstances that they live in. And while meant to ensure inclusion, these systems, in effect, often exclude many of the people who need or are entitled to such services.

21:30 And Jordan is just one example among many, but there are several reports about this, some of which are linked in the reading material here. So looking to the policies, though, in Jordan, while I won't spend too much time on this, because the focus here really is on, kind of, the external support and funding for it. In terms of policies, Jordan itself is often considered to be ahead of its peers in putting together an AI strategy with five pillars. And one of these focuses on applying AI tools to improve the efficiency of government services and the public sector. And the country has also put together an AI ethics charter to govern the use of AI domestically.

22:15 As I mentioned, in addition to policy and regulation at a domestic level, there are also other avenues where parameters for the use of AI and data can be identified or set. For example, there are many diplomatic efforts that are taking place right now, we see that a big focus of the G20 in 2023 was the idea of harnessing artificial intelligence responsibly for good and for all. The Global Partnership on AI has, you know, been working towards increasing their output, and they have a recent Ministerial Declaration of 2023. There was an AI Safety Summit that the UK hosted this past year. And we also have as an older example of such efforts, the OECD Principles on AI, which have inspired a lot of different governments' policies and frameworks on AI principles and ethics. But while these typically talk about how, you know, developing AI responsibly and for the public interest is very important, they're not very specifically targeted towards government use and its regulation itself, and they look more broadly to the use of AI and development of AI across public and private sectors.

23:44 Another example of, you know, actions outside of policy and regulation is the way courts look at AI. They typically look at both the potential for use of automated systems to support government and law enforcement actions, and the need to limit such use in the context of human rights violations. For example, we see famously a court order in the Netherlands, which ordered the government to stop the use of its digital fraud detection tool, which targeted, you know, people in poorer neighborhoods who are receiving welfare benefits for the purpose of detecting welfare fraud. But the court found that this violated the human rights of these individuals and asked the government to stop using the system. In India, another example is a discussion within courts where there was, you know, social media companies highlighted the possibility of using AI for proactive detect detection of child sexual abuse and abuse of material before the Supreme Court of India. There's many other examples of these kinds of situations where courts are asking private companies to work on these issues towards enforcement of public law. In the US, we also see multiple issues with the use of automated tools for welfare distribution services, and so these issues exist across many countries. And, of course, the other way in which AI and courts come together could be in the incorporation of AI into judicial or administrative processes in the courts. And that could be seen as one more way in which AI is used in the public sector.

25:37 So while we've discussed the many different use cases and examples for AI and GovTech, as well as some of the policy and regulatory mechanisms that could be applied here, we see that, you know, usually this is not enough, because there's several potential problems that could arise from this kind of use. As we've seen, just in these examples that were discussed, there could be discriminatory or faulty systems. This is particularly common in the case of welfare benefits or law enforcement uses of AI. But also, in addition to the context of, you know, individual rights and service delivery, there could be problems in structuring administrative offices better and industrial processes on these subjects as well. So we look at discriminatory and faulty systems as one of the problems. Another big problem that exists is the issue of privacy and cybersecurity, and kind of the structuring of these technical systems to work effectively, while also using the personal data of the individuals it serves, properly and also securely.

26:58 We also see that there could be risks to the principles of, you know, accountability, transparency and oversight that we expect our governments to follow in this context. This could be because of the use of black box systems. In other places, it could be because there's limited technical capacity in the relevant government institutions. Often we see that many of these systems result in some sort of function or mission creep. So while they're originally put in place for a targeted use, these systems are increasingly used for other purposes, making it harder to use the data in a proper manner as envisioned. We also see that there are potential problems when we use datasets or AI, that is trained by private actors, such as the vendors and consultants who are usually involved in procurement processes by governments, because this, again, increases risks in that process. And then the last, you know, issue that we see that comes up is the recent developments in the context of generative AI, because we still need to understand how this may be used and what the risks are. And while governments are pushing to identify use cases for generative AI within their functions, a lot of folks argue that it is a technology that is ill equipped in this context.

28:39 So we see also from this particular example, that there is a heightened risk of, you know, there's a presumption of accuracy and this potential that technology will solve the problems. And AI, particularly, will solve many of the problems that the government faces in the administration of their functions and services, which is also something we need to be cautious about. And as an underlying issue here, we see that many countries are now treating GovTech as infrastructure. That is they are building technology, both hardware and software, that will act as the base, or the infrastructure that underlies tech based products and services that are built in the future. And this could be both private and public sector technologies. And with the push to integrate AI into these systems, we need to ensure that, if we are doing so, it's done in the right way with adequate impact assessments, oversight, etc. before we bake bad technological systems into this infrastructure and make it more difficult to undo.

29:53 So, that's about it from me today. I hope that you were able to get some insight into how governments are using AI and how they're thinking about potential use cases for AI. But also some of the concerns that we need to be wary of and that, you know, we need to make sure are incorporated into the policies and frameworks that are coming out on this subject. Thank you for listening.